

MANAN SHARMA

MCA (AI & ML) — MIT-WPU Pune

✉ manansharmajpr@gmail.com | 📞 +91-9828932000 | 📍 Pune, Maharashtra

🌐 linkedin.com/in/manan-sharma-780b13254 | 🐙 github.com/MANANSHARMA275 | 🌐 manansharma275.github.io

Objective

MCA (AI & ML) student with hands-on experience building data pipelines, relational schemas, and ML-driven analytics in production — comfortable across Python, SQL, and cloud-deployed systems, with a track record of shipping and testing real data infrastructure end to end. Seeking data analytics / data engineering roles to build toward a long-term goal in AI/ML engineering.

Education

MIT World Peace University

Master of Computer Applications (MCA) — AI & ML Electives | CGPA: 7.60

Pune, Maharashtra

08/2025 – 05/2027

ICFAI University

Bachelor of Computer Applications — AI & IoT | Percentage: 74.50%

Jaipur, Rajasthan

07/2022 – 05/2025

Technical Skills

Data & Languages: Python · SQL · Pandas · NumPy · Statistics

Databases: PostgreSQL · SQLAlchemy · Flask-Migrate

Backend & APIs: FastAPI · Flask · REST API Development · Celery · Redis · Pydantic

Cloud, DevOps & Testing: Docker · Docker Compose · AWS (EC2, S3, RDS) · Git · GitHub · Pytest

ML & GenAI: CatBoost · Streamlit · GenAI / RAG Basics

Experience

MineralLaw.in — Mining Compliance Platform

Pune, India

Co-Founder — Data & Backend Engineer 🧑‍💻

2026 – Present

- Designed and built an automated compliance data pipeline (Celery Beat) that processes statutory deadline events for mining lease holders, with fingerprint-based idempotency guaranteeing exactly-once processing across repeated pipeline runs.
- Modeled the relational schema in PostgreSQL for compliance records and delivery logs via SQLAlchemy, with schema changes version-controlled through Flask-Migrate to support iterative evolution.
- Built a rule-based deadline calculator driven by a central configuration of statutory rules, converting regulatory dates into structured, queryable event data; delivery tasks guarded by consent checks, a kill switch, and database-level duplicate constraints.
- Applied data-minimisation practices (masked phone numbers in delivery records) and built a paginated, filterable admin log for auditing reminder deliveries.
- Containerized the stack with Docker Compose and deployed on AWS (EC2, RDS); validated with a 301-test Pytest suite across unit and integration layers.

Projects

Food Delivery Time Prediction | Python, Pandas, CatBoost, Streamlit 🧑‍💻

- Performed EDA and feature engineering on 45,593 delivery orders: engineered geo features (haversine and Manhattan-proxy distance, distance buckets) and repaired invalid coordinates via city-wise median imputation.
- Trained and benchmarked CatBoost against linear regression and Random Forest baselines; CatBoost reached MAE 5.70 min and RMSE 7.26 (R^2 0.40) on a held-out 20% split — a 7.3% RMSE improvement over linear regression and 4.4% over Random Forest.
- Deployed as a Streamlit app performing live ETA prediction with input and coordinate validation.

Tokenmeter — LLM Inference Cost Calculator | JavaScript, HTML, CSS 🌐 🧑‍💻

2026

- Built a zero-dependency single-page tool computing monthly LLM cost across six model tiers with live recomputation and USD/INR localisation, pricing defaults verified against provider pages.

Certifications

Machine Learning with Python | IBM / Coursera

October 2024

Mastering Data Structures & Algorithms using C and C++ | Udemy — Abdul Bari

August 2026 – Present

Achievements & Activities

Identified the MineralLaw.in opportunity independently — existing compliance portals serve regulators and large operators, not small lease holders — then defined the product scope and built it from an empty repository.

Built and presented working prototypes in two-person teams under time-boxed hackathon conditions, including YOLO Pune 2026 at MIT-WPU.